

# ML-7

## PROJECT 1

Muhammad Naqeeb | Machine-Learning / project 1 | August 18,  
2021

# What is structured data?



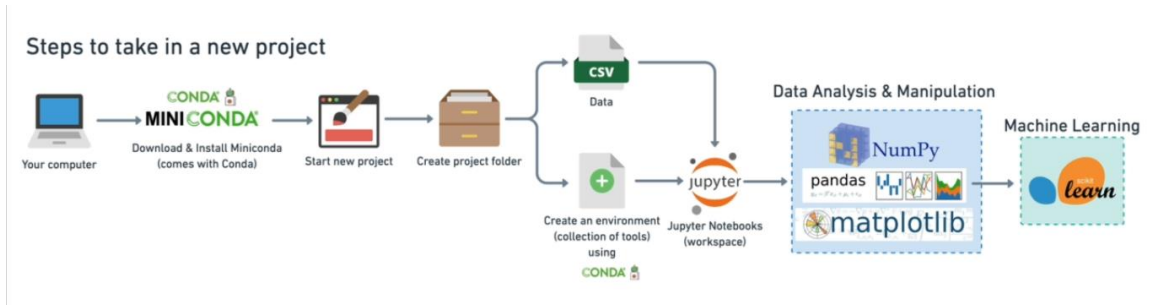
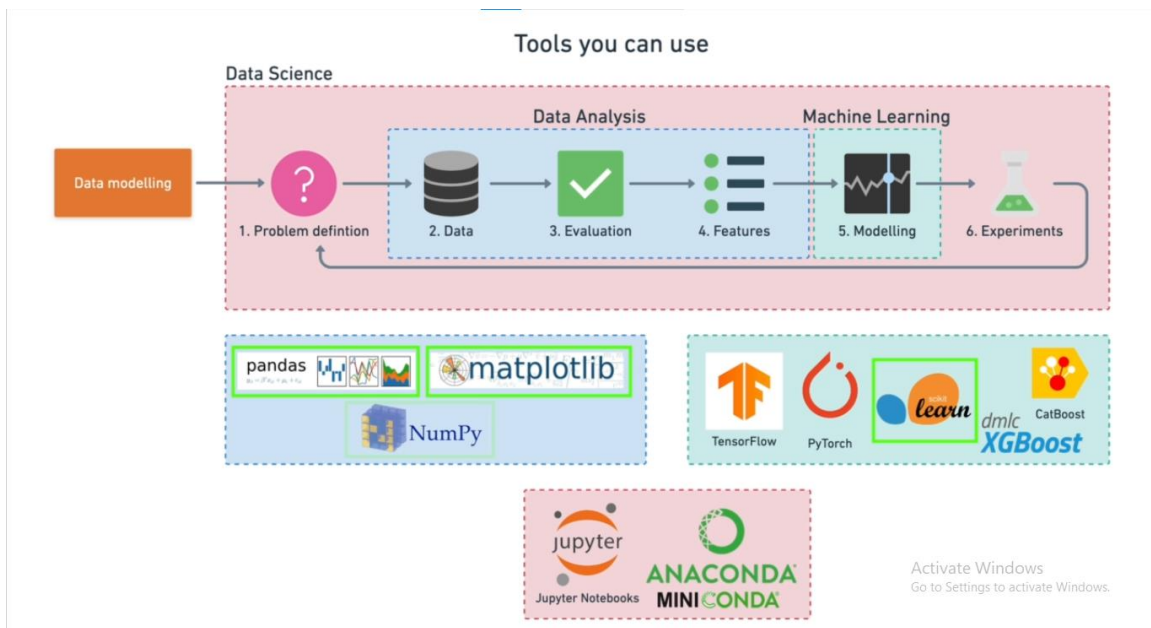
Feature variables					Target
ID	Weight	Sex	Heart Rate	Chest pain	Heart disease?
4528	110kg	M	81	4	Yes
5681	64kg	F	61	1	No
7911	81kg	M	57	0	No

Table 1.0: Patient records

Activate Windows

Steps in a full machine learning project





# Predicting heart disease using machine learning

This notebook looks into using various Python based machine learning and data science libraries which are our tools in an attempt to build a machine learning model capable of predicting whether or not someone has heart disease based on their medical attributes.

**We're going to take the following approach.**

1. Problem definition
2. Data
3. Evaluation
4. Features
5. Modelling
6. Experimentation

## 1. Problem definition

In a statement,

Given clinical parameters about a patient, can we predict whether or not they have heart disease"?

## 2. Data

The original data came from the Cleveland database from UCI machine learning repository There is also a version of it available on Kaggle.

## 3. Evaluation

If we can reach 95 percent accuracy at predicting whether or not a patient has heart disease during the proof of concept, pursue the project.

## 4. Features

This is where you'll get different information about each of the features in your data.

## Preparing the tools.

We're going to use Panda's matplotlib and Numpy for data analysis and manipulation.

```
# Import all the tools we need

# Regular EDA (exporatory data analysis) and plotting libraries

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

#we want our plots to appear inside the notebook

%matplotlib inline

#Model from Scikit-learn
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier

# Model Evaluation
from sklearn.model_selection import train_test_split,
cross_val_score
from sklearn.model_selection import RandomizedSearchCV,
GridSearchCV
from sklearn.metrics import confusion_matrix,
classification_report
from sklearn.metrics import precision_score, recall_score,
f1_score
from sklearn.metrics import plot_roc_curve
```

## Load data

```
df = pd.read_csv("11.2 heart-disease.csv")
df
```

```
In [10]: df = pd.read_csv("11.2 heart-disease.csv")
df
```

Out[10]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
298	57	0	0	140	241	0	1	123	1	0.2	1	0	3	0
299	45	1	3	110	264	0	1	132	0	1.2	1	0	3	0
300	68	1	0	144	193	1	1	141	0	3.4	1	2	3	0
301	57	1	0	130	131	0	1	115	1	1.2	1	1	3	0
302	57	0	1	130	236	0	0	174	0	0.0	1	1	2	0

303 rows × 14 columns

## Data Explortion (exploratory data analysis or EDA)

So the goal here is to find out more about the data and become a subject matter expert on the data set you're working with

1. what questions are you trying to solve.
2. what kind of data do we have. And how do we treat different types?
3. What's missing from the data. And how do you deal with it?
4. where are the outliers and why should you care about them?
5. How can you add, change or remove features to get more out of your data.

```
# let's find out how many of each class there
df["target"].value_counts()
```


```
In [16]: # let's find out how many of each class there
df["target"].value_counts()
```

Out[16]:

1	165
0	138

Name: target, dtype: int64

```
df["target"].value_counts().plot(kind="bar", color=["salmon",
"lightblue"] )
```



```
df.shape
```

```
df.info()
```

```
In [14]: df.shape
```

```
Out[14]: (303, 14)
```

```
In [18]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 303 entries, 0 to 302  
Data columns (total 14 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   age         303 non-null   int64  
1   sex         303 non-null   int64  
2   cp          303 non-null   int64  
3   trestbps    303 non-null   int64  
4   chol        303 non-null   int64  
5   fbs         303 non-null   int64  
6   restecg     303 non-null   int64  
7   thalach     303 non-null   int64  
8   exang       303 non-null   int64  
9   oldpeak     303 non-null   float64  
10  slope       303 non-null   int64  
11  ca          303 non-null   int64  
12  thal        303 non-null   int64  
13  target      303 non-null   int64  
dtypes: float64(1), int64(13)  
memory usage: 33.3 KB
```

```
df.isna().sum() # checking missing
```

```
In [20]: df.isna().sum() # checking missing
```

```
Out[20]: age         0  
sex         0  
cp          0  
trestbps    0  
chol        0  
fbs         0  
restecg     0  
thalach     0  
exang       0  
oldpeak     0  
slope       0  
ca          0  
thal        0  
target      0  
dtype: int64
```

```
df.describe()
```



```
In [21]: df.describe()
```

Out[21]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	0.528053	149.646865	0.326733	1.039604	1.399340	0.729373
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	0.525860	22.905161	0.469794	1.161075	0.616226	1.022606
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000	133.500000	0.000000	0.000000	1.000000	0.000000
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000	153.000000	0.000000	0.800000	1.000000	0.000000
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	1.000000	166.000000	1.000000	1.600000	2.000000	1.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.000000	1.000000	6.200000	2.000000	4.000000

## Heart Disease Frequency according to Sex

```
df.sex.value_counts()
```

```
# compare target column with sex column
pd.crosstab(df.target, df.sex)
```

```
In [23]: df.sex.value_counts()

Out[23]: 1    207
         0     96
         Name: sex, dtype: int64
```


```
In [25]: # compare target column with sex column
pd.crosstab(df.target, df.sex)
```

```
Out[25]:
      sex  0    1
target
0      24  114
1      72   93
```

```
# Create a plot of crosstab
pd.crosstab(df.target, df.sex).plot(kind="bar",
                                     figsize=(10,6),
                                     color=["salmon", "lightblue"])

plt.title("Heart Disease Frequency for sex")
```

```
plt.xlabel("0 = No Disease, 1 = Disease")
plt.ylabel("Amount")
plt.legend(["Female", "Male"])
plt.xticks(rotation = 0);
```




## Age vs. Max Heart Rate for Heart Disease

```
# Create another figure
plt.figure(figsize=(10, 6))


# Scatter with positive example
plt.scatter(df.age[df.target==1],
            df.thalach[df.target==1],
            c="salmon")

# Scatter with negative example
plt.scatter(df.age[df.target==0],
            df.thalach[df.target==0],
            c="lightblue")
```

```
# Add some helpful info
plt.title("Heart Disease in function of Age and Max Heart Rate")
plt.xlabel("Age")
plt.ylabel("Max Heart Rate")
plt.legend(["Disease", "No Disease"])
```



```
# Check the distribution of the age column with a histogram
df.age.plot.hist();
```



## Heart Disease Frequency per Chest pain Type

```
pd.crosstab(df.cp, df.target)
```

```
In [42]: pd.crosstab(df.cp, df.target)
```


```
Out[42]:
```

target	0	1
cp		
0	104	39
1	9	41
2	18	69
3	7	16

```
# Make the crosstab more visual
# Create a plot of crostab
pd.crosstab(df.cp, df.target).plot(kind="bar",
                                     figsize=(10,6),
                                     color=["salmon", "lightblue"])

plt.title("Heart Disease Frquency per Chest pain Type")
plt.xlabel("Chest Pain Typr")
plt.ylabel("Amount")
```

```
plt.legend(["No disease", "disease"])
plt.xticks(rotation = 0);
```



```
# Make a correlation matrix
df.corr()
```


```
In [44]: # Make a correlation matrix
df.corr()
```

Out[44]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
age	1.000000	-0.098447	-0.068653	0.279351	0.213678	0.121308	-0.116211	-0.398522	0.096801	0.210013	-0.168814	0.276326	0.068001	-0.225439
sex	-0.098447	1.000000	-0.049353	-0.056769	-0.197912	0.045032	-0.058196	-0.044020	0.141664	0.096093	-0.030711	0.118261	0.210041	-0.280937
cp	-0.068653	-0.049353	1.000000	0.047608	-0.076904	0.094444	0.044421	0.295762	-0.394280	-0.149230	0.119717	-0.181053	-0.161736	0.433798
trestbps	0.279351	-0.056769	0.047608	1.000000	0.123174	0.177531	-0.114103	-0.046698	0.067616	0.193216	-0.121475	0.101389	0.062210	-0.144931
chol	0.213678	-0.197912	-0.076904	0.123174	1.000000	0.013294	-0.151040	-0.009940	0.067023	0.053952	-0.004038	0.070511	0.098803	-0.085239
fbs	0.121308	0.045032	0.094444	0.177531	0.013294	1.000000	-0.084189	-0.008567	0.025665	0.005747	-0.059894	0.137979	-0.032019	-0.028046
restecg	-0.116211	-0.058196	0.044421	-0.114103	-0.151040	-0.084189	1.000000	0.044123	-0.070733	-0.058770	0.093045	-0.072042	-0.011981	0.137230
thalach	-0.398522	-0.044020	0.295762	-0.046698	-0.009940	-0.008567	0.044123	1.000000	-0.378812	-0.344187	0.386784	-0.213177	-0.096439	0.421741
exang	0.096801	0.141664	-0.394280	0.067616	0.067023	0.025665	-0.070733	-0.378812	1.000000	0.288223	-0.257748	0.115739	0.206754	-0.436757
oldpeak	0.210013	0.096093	-0.149230	0.193216	0.053952	0.005747	-0.058770	-0.344187	0.288223	1.000000	-0.577537	0.222682	0.210244	-0.430696
slope	-0.168814	-0.030711	0.119717	-0.121475	-0.004038	-0.059894	0.093045	0.386784	-0.257748	-0.577537	1.000000	-0.080155	-0.104764	0.345877
ca	0.276326	0.118261	-0.181053	0.101389	0.070511	0.137979	-0.072042	-0.213177	0.115739	0.222682	-0.080155	1.000000	0.151832	-0.391724
thal	0.068001	0.210041	-0.161736	0.062210	0.098803	-0.032019	-0.011981	-0.096439	0.206754	0.210244	-0.104764	0.151832	1.000000	-0.344029
target	-0.225439	-0.280937	0.433798	-0.144931	-0.085239	-0.028046	0.137230	0.421741	-0.436757	-0.430696	0.345877	-0.391724	-0.344029	1.000000

```
# Let's make our correlation matrix a little prettier
```

```
corr_matrix = df.corr()
fig, ax = plt.subplots(figsize = (15,10))
ax =sns.heatmap(corr_matrix,
                 annot=True,
                 linewidth=0.5,
                 fmt=".2f",
                 cmap="YlGnBu")
```



A higher positive value means a potential positive correlation and a higher negative value means a potential negative correlation or a decrease. this is saying is as C.P. goes up the target value also increases

**Correction: Negative correlation** = a relationship between two variables in which one variable increases as the other decreases

## 6. Modelling

```
# Split data into x and y
X = df.drop("target", axis = 1)
y=df["target"]
```

```
# split data into train and test sets
np.random.seed(42)

# # split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2)
```

Now we've got our data split into training and test sets, it's time to build a machine learning model.

We'll train it (find the patterns) on the training set.

And we'll test it (use the patents) on the test set.

We're going to try 3 different machine learning models:

1. Logistic Regression
2. K-Nearest Neighbours Classifier
3. Random Forest Classifier

```
# Put models in a dictionary
models = {"Logistic Regression":LogisticRegression(),
          "KNN":KNeighborsClassifier(),
          "Random Forest":RandomForestClassifier()}

# Create a function to fit and score models
def fit_and_score(models,X_train,X_test, y_train, y_test):

    """
    Fits and evaluates given machine learning models.
    models : a dict of different scikit-Learn machine learning
models
    X_train : training data (no labels)
```

```

X_test : trsting data (no labels)
y_train : training labels
y_test : test labels
"""

# Set random seed
np.random.seed(42)
# Make a dictionary to keep model score
model_scores={}

# Loop through models
for name, model in models.items():
    # Fit the model to data
    model.fit(X_train, y_train)
    # Evaluate the mode and append its score to model_score
    model_scores[name] = model.score(X_test, y_test)

return model_scores

```

```

# let's see how each model perform
model_scores = fit_and_score(models = models,
                              X_train = X_train
                              , X_test = X_test
                              , y_train = y_train
                              , y_test = y_test)

model_scores

```



```

In [24]: # Let's see how each model perform
model_scores = fit_and_score(models = models,
                              X_train = X_train
                              , X_test = X_test
                              , y_train = y_train
                              , y_test = y_test)

model_scores

C:\Users\toshiba c55t-a\desktop\sample_project\env\lib\site-packages\sk
learn\linear_model\_logistic.py:763: ConvergenceWarning: lbfgs failed t
o converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown
in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-
regression
n_iter_i = _check_optimize_result(

Out[24]: {'Logistic Regression': 0.8852459016393442,
          'KNN': 0.6885245901639344,
          'Random Forest': 0.8360655737704918}


```

## Model Comparison

```

model_compare = pd.DataFrame(model_scores, index=["accuracy"])
model_compare.T.plot.bar()

```



Now we've got a baseline model... and we know a models first predictions aren't always what we should base our next steps off.

## What should we do?

let's look at the following

- Hyperparameter tuning
- Feature importance
- Confusion matrix
- Cross-validation
- Precision
- Recall
- F1 score
- Classification report
- ROC curve
- Area under the curve (AUC)

# Classification and Regression metrics

Classification	Regression
Accuracy	R <sup>2</sup> (r-squared)
Precision	Mean absolute error (MAE)
Recall	Mean squared error (MSE)
<b>F1</b>	Root mean squared error (RMSE)

**Bold** = default evaluation in Scikit-Learn

Activate Windows  
Go to Settings to activate Wind

## HYPERPARAMETER TUNING

```
# Let's tune KNN

train_scores = []
test_scores = []

# Create a list of different values for n_neighbors
neighbors = range(1,21)

# Setup KNN instance
knn = KNeighborsClassifier()

# loop through different n-neighbors
for i in neighbors:
    knn.set_params(n_neighbors = i)

    # Fit the algorithms
    knn.fit(X_train, y_train)

    # Update the training scores list
    train_scores.append(knn.score(X_train, y_train))
```

```
# Update the test scores list
test_scores.append(knn.score(X_test, y_test))
```

```
train_scores
```


```
In [32]: train_scores
Out[32]: [1.0,
          0.8099173553719008,
          0.7727272727272727,
          0.743801652892562,
          0.7603305785123967,
          0.7520661157024794,
          0.743801652892562,
          0.7231404958677686,
          0.71900826446281,
          0.6942148760330579,
          0.7272727272727273,
          0.6983471074380165,
          0.6900826446280992,
          0.6942148760330579,
          0.6859504132231405,
          0.6735537190082644,
          0.6859504132231405,
          0.6652892561983471,
          0.6818181818181818,
          0.6694214876033058]
```

```
test_scores
```

```
In [33]: test_scores
```

```
Out[33]: [0.6229508196721312,  
          0.639344262295082,  
          0.6557377049180327,  
          0.6721311475409836,  
          0.6885245901639344,  
          0.7213114754098361,  
          0.7049180327868853,  
          0.6885245901639344,  
          0.6885245901639344,  
          0.7049180327868853,  
          0.7540983606557377,  
          0.7377049180327869,  
          0.7377049180327869,  
          0.7377049180327869,  
          0.6885245901639344,  
          0.7213114754098361,  
          0.6885245901639344,  
          0.6885245901639344,  
          0.7049180327868853,  
          0.6557377049180327]
```

```
plt.plot(neighbors, train_scores, label="Train Score")  
  
plt.plot(neighbors, test_scores, label="test Score")  
plt.xticks(np.arange(1, 21, 1))  
plt.xlabel("Number of neighbors")  
plt.ylabel("Model score")  
plt.legend()  
  
print(f"Maximum KNN score on the test data:  
{max(test_scores)*100:.2f}%")
```



## HYPERPARAMETER TUNING WITH RANDOMIZEDSEARCHCV

We're going to tune

- LogisticRegression()
- RandomForestClassifier()

... using RandomizedSearchCV

```
# create a hyper parameter grid for logisticRegression()

log_reg_grid = {"C": np.logspace(-4, 4, 20),
                "solver": ["liblinear"]}

# Create a hyperparameter grid for RandomForestClassifier
rf_grid = {"n_estimators": np.arange(10, 1000, 50),
           "max_depth": [None, 3, 5, 10],
           "min_samples_split": np.arange(2, 20, 2),
           "min_samples_leaf": np.arange(1, 20, 2)}
```

Now we've got hyperparameter grid setup for each of our models, let's tune them using RandomizedSearchCV

```
# tune logisticRegression

np.random.seed(42)
# Setup random hyperparameter search for logisticRegression
rs_log_reg = RandomizedSearchCV(LogisticRegression(),
                                param_distributions=log_reg_grid,
                                cv=5,
                                n_iter = 20,
                                verbose = True)

# fit random hyperparameter search for logisticRegression
rs_log_reg.fit(X_train, y_train)
```

```
In [42]: # tune logisticRegression

np.random.seed(42)
# Setup random hyperparameter search for logisticRegression
rs_log_reg = RandomizedSearchCV(LogisticRegression(),
                                param_distributions=log_reg_grid,
                                cv=5,
                                n_iter = 20,
                                verbose = True)

# fit random hyperparameter search for logisticRegression
rs_log_reg.fit(X_train, y_train)

Fitting 5 folds for each of 20 candidates, totalling 100 fits
Out[42]: RandomizedSearchCV(cv=5, estimator=LogisticRegression(), n_iter=20,
                             param_distributions={'C': array([1.00000000e-04, 2.6
3665090e-04, 6.95192796e-04, 1.83298071e-03,
4.83293024e-03, 1.27427499e-02, 3.35981829e-02, 8.85866790e-02,
2.33572147e-01, 6.15848211e-01, 1.62377674e+00, 4.28133240e+00,
1.12883789e+01, 2.97635144e+01, 7.84759970e+01, 2.06913808e+02,
5.45559478e+02, 1.43844989e+03, 3.79269019e+03, 1.00000000e+0
4]),
                             'solver': ['liblinear']}),
                             verbose=True)
```

```
rs_log_reg.best_params_
```

```
rs_log_reg.score(X_test, y_test)
```

```
In [43]: rs_log_reg.best_params_
```

```
Out[43]: {'solver': 'liblinear', 'C': 0.23357214690901212}
```

```
In [45]: rs_log_reg.score(X_test, y_test)
```

```
Out[45]: 0.9344262295081968
```

Now we've tuned `logisticRegression()`, let's do the same for `RandomForestClassifier()`.....

```
np.random.seed(42)
# Setup random hyperparameter search for RandomForestClassifier()
rs_rf = RandomizedSearchCV(RandomForestClassifier(),
                           param_distributions=rf_grid,
                           cv=5,
                           n_iter = 20,
                           verbose = True)

# fit random hyperparameter search for RandomForestClassifier()
rs_rf.fit(X_train, y_train)
```



```

In [48]: > np.random.seed(42)
# Setup random hyperparameter search for RandomForestClassifier()
rs_rf = RandomizedSearchCV(RandomForestClassifier(),
                           param_distributions=rf_grid,
                           cv=5,
                           n_iter = 20,
                           verbose = True)

# fit random hyperparameter search for RandomForestClassifier()
rs_rf.fit(X_train, y_train)

Fitting 5 folds for each of 20 candidates, totalling 100 fits
Out[48]: RandomizedSearchCV(cv=5, estimator=RandomForestClassifier(), n_iter=20,
                           param_distributions={'max_depth': [None, 3, 5, 10],
                                                'min_samples_leaf': array([ 1,
3,  5,  7,  9, 11, 13, 15, 17, 19]),
                                                'min_samples_split': array([ 2,
4,  6,  8, 10, 12, 14, 16, 18]),
                                                'n_estimators': array([ 10,  6
0, 110, 160, 210, 260, 310, 360, 410, 460, 510, 560, 610,
660, 710, 760, 810, 860, 910, 960])},
                           verbose=True)

```

```

# Find the best hyperparameters
rs_rf.best_params_

```

```

# Evaluate the randomized search RandomForestClassifier model
rs_rf.score(X_test, y_test)

```

```

In [49]: > # Find the best hyperparameters
rs_rf.best_params_

```

```

Out[49]: {'n_estimators': 210,
          'min_samples_split': 4,
          'min_samples_leaf': 19,
          'max_depth': 3}

```

```

In [50]: > # Evaluate the randomized search RandomForestClassifier model
rs_rf.score(X_test, y_test)

```

```

Out[50]: 0.8688524590163934

```

## HYPERPARAMETERS TUNING WITH GRIDSEARCHCV

Since our logisticRegression model provides the best scores so far, we'll try and improve them again using GridSearchCV

```
# different hyperparameters for our LogisticRegression model
log_reg_grid = {"C": np.logspace(-4, 4, 30),
                "solver": ["liblinear"]}

# Setup grid hyperparameter search for LogisticRegression model
gs_log_reg = GridSearchCV(
    LogisticRegression(),
    param_grid=log_reg_grid,
    cv=5,
    verbose = True
)

# Fit grid hyperparameter search model
gs_log_reg.fit(X_train, y_train)
```

```
In [53]: # different hyperparameters for our LogisticRegression model
log_reg_grid = {"C": np.logspace(-4, 4, 30),
                "solver": ["liblinear"]}

# Setup grid hyperparameter search for LogisticRegression model
gs_log_reg = GridSearchCV(
    LogisticRegression(),
    param_grid=log_reg_grid,
    cv=5,
    verbose = True
)

# Fit grid hyperparameter search model
gs_log_reg.fit(X_train, y_train)
```

Fitting 5 folds for each of 30 candidates, totalling 150 fits

```
Out[53]: GridSearchCV(cv=5, estimator=LogisticRegression(),
                    param_grid={'C': array([1.00000000e-04, 1.88739182e-04, 3.
56224789e-04, 6.72335754e-04,
1.26896100e-03, 2.39502662e-03, 4.52035366e-03, 8.53167852e-03,
1.61026203e-02, 3.03919538e-02, 5.73615251e-02, 1.08263673e-01,
2.04335972e-01, 3.85662042e-01, 7.27895384e-01, 1.37382380e+00,
2.59294380e+00, 4.89390092e+00, 9.23670857e+00, 1.74332882e+01,
3.29034456e+01, 6.21016942e+01, 1.17210230e+02, 2.21221629e+02,
4.17531894e+02, 7.88046282e+02, 1.48735211e+03, 2.80721620e+03,
5.29831691e+03, 1.00000000e+04]),
                    'solver': ['liblinear']},
                    verbose=True)
```

Activate Windows

```
# best hyperparameters
gs_log_reg.best_params_
```

```
# Evaluate the grid search LogisticRegression model
gs_log_reg.score(X_test, y_test)
```

```
In [54]: ► # best hyperparameters
          gs_log_reg.best_params_
```

```
Out[54]: {'C': 0.20433597178569418, 'solver': 'liblinear'}
```

```
In [55]: ► # Evaluate the grid search LogisticRegression model
          gs_log_reg.score(X_test, y_test)
```

```
Out[55]: 0.8852459016393442
```

## Evaluating our tuned machine learning classifier, beyond accuracy

- ROC curve and AUC score
- Confusion matrix
- Classification report
- Precision
- Recall
- F1 score

.... and it would be great if cross-validation was used where possible.

To make comparisons and evaluate our trained model, first we need to make predictions.

```
# Make predictions with tuned model
y_preds = gs_log_reg.predict(X_test)
```

```
y_preds
```


```
y_test
```

```
In [56]: # Make predictions with tuned model
y_preds = gs_log_reg.predict(X_test)
```

```
In [57]: y_preds
```

```
Out[57]: array([0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1,
0,
           0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1,
1,
           1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0], dtype=int64)
```

```
# Plot ROC curve and calculate and caculate AUC metric
plot_roc_curve(gs_log_reg, X_test, y_test)
```



```
# Confution matrix
```

```
print(confusion_matrix(y_test, y_preds))
```

```
In [60]: # Confution matrix
print(confusion_matrix(y_test, y_preds))

[[25  4]
 [ 3 29]]
```


```
sns.set(font_scale=1.5)

def plot_conf_mat(y_test, y_preds):
    """
    Plots a nice looking confusion metrix using Seaborn's
    heatmap()
    """

    fig, ax = plt.subplots(figsize=(3,3))
    ax = sns.heatmap(confusion_matrix(y_test, y_preds),
                      annot=True,
                      cbar=False )

    plt.xlabel("True Label")
    plt.ylabel("Predicted Label")

plot_conf_mat(y_test, y_preds)
```



Now we've got a ROC curve and AUC metric and a confusion matrix, Let's get a classification report as well as cross validated precision, recall and fi-score

```
print(classification_report(y_test, y_preds))
```

```
In [67]: print(classification_report(y_test, y_preds))
```

	precision	recall	f1-score	support
0	0.89	0.86	0.88	29
1	0.88	0.91	0.89	32
accuracy			0.89	61
macro avg	0.89	0.88	0.88	61
weighted avg	0.89	0.89	0.89	61

## Calculate evaluation metrics using cross validation

we're going to calculate accuracy, precision, recall and fi score of our model using cross-validation, and to do so we'll be using `cross_val_score()`

```
# Check best hyperparameters
gs_log_reg.best_params_
```

```
# Create a new classifier with best parameters
clf = LogisticRegression(C=0.20433597178569418,
                        solver= "liblinear")
```

```
# Cross-validated accuracy
cv_acc = cross_val_score(clf,
                        X,
                        Y,
                        cv=5,
                        scoring = "accuracy")

cv_acc
```

```
cv_acc = np.mean(cv_acc)
cv_acc
```

```
In [63]: # Check best hyperparameters
gs_log_reg.best_params_
```

```
Out[63]: {'C': 0.20433597178569418, 'solver': 'liblinear'}
```

```
In [65]: # Create a new classifier with best parameters
clf = LogisticRegression(C=0.20433597178569418,
                        solver= "liblinear")
```

```
In [70]: # Cross-validated accuracy
cv_acc = cross_val_score(clf,
                        X,
                        Y,
                        cv=5,
                        scoring = "accuracy")

cv_acc
```

```
Out[70]: array([0.81967213, 0.90163934, 0.86885246, 0.88333333, 0.75      ])
```

```
In [72]: cv_acc = np.mean(cv_acc)
cv_acc
```

```
Out[72]: 0.8446994535519124
```

Activate Windows

Go to Settings to activate Windows.

```
# Cross-validated precision
cv_precision = cross_val_score(clf,
                                X,
                                y,
                                cv=5,
                                scoring = "precision")
cv_precision = np.mean(cv_precision)
cv_precision
```

```
# Cross-validated recall
cv_recall = cross_val_score(clf,
                             X,
                             y,
                             cv=5,
                             scoring = "recall")
cv_recall = np.mean(cv_recall)
cv_recall
```

```
# Cross-validated f1-score
cv_f1 = cross_val_score(clf,
                         X,
                         y,
                         cv=5,
                         scoring = "f1")
cv_f1 = np.mean(cv_f1)
cv_f1
```



```
In [73]: # Cross-validated precision
cv_precision = cross_val_score(clf,
                                X,
                                y,
                                cv=5,
                                scoring = "precision")
cv_precision = np.mean(cv_precision)
cv_precision
```

Out[73]: 0.8207936507936507

```
In [74]: # Cross-validated recall
cv_recall = cross_val_score(clf,
                             X,
                             y,
                             cv=5,
                             scoring = "recall")
cv_recall = np.mean(cv_recall)
cv_recall
```


Out[74]: 0.9212121212121213

```
In [75]: # Cross-validated f1-score
cv_f1 = cross_val_score(clf,
                         X,
                         y,
                         cv=5,
                         scoring = "f1")
cv_f1 = np.mean(cv_f1)
cv_f1
```

Out[75]: 0.8673007976269721

```
# Visualize cross-validation metrics
cv_metrics = pd.DataFrame({"Accuracy":cv_acc,
                           "Precision":cv_precision,
                           "Recall":cv_recall,
                           "F1":cv_f1},
                           index=[0])

cv_metrics.T.plot.bar(title="cross-validation classification
metrics",
legend = False)
```



## Feature Importance

Feature Importance is another as asking, "which features contributed most to the outcomes of the model and how did they contribute?"

finding feature importance is different for each machine learning model. One way to find feature importance is to search for **(MODEL NAME) feature importance**

let's find the feature important for now logisticRegression model...

```
# Fit an instance of logisticRegression

clf = LogisticRegression(C=0.20433597178569418,solver =
"liblinear")

clf.fit(X_train, y_train)
```

```
# check coef_  
clf.coef_
```

```
In [83]: # Fit an instance of LogisticRegression  
  
clf = LogisticRegression(C=0.20433597178569418,solver = "liblinear")  
clf.fit(X_train, y_train)  
  
Out[83]: LogisticRegression(C=0.20433597178569418, solver='liblinear')
```

```
In [84]: # check coef_  
clf.coef_  
  
Out[84]: array([[ 0.00316728, -0.86044652,  0.6606704 , -0.01156993, -0.0016637  
5,  
               0.04386107,  0.31275848,  0.02459362, -0.60413081, -0.5686280  
3,  
               0.45051628, -0.63609898, -0.67663373]])
```

```
# Match coef's of features to columns  
feature_dict = dict(zip(df.columns, list(clf.coef_[0])))  
feature_dict
```

```
In [85]: # Match coef's of features to columns  
feature_dict = dict(zip(df.columns, list(clf.coef_[0])))  
feature_dict  
  
Out[85]: {'age': 0.0031672806268220445,  
'sex': -0.8604465226286001,  
'cp': 0.6606703996492814,  
'trestbps': -0.011569930743501303,  
'chol': -0.001663745833540806,  
'fbs': 0.043861067871676124,  
'restecg': 0.3127584791782968,  
'thalach': 0.02459361509185037,  
'exang': -0.6041308102637141,  
'oldpeak': -0.5686280255489925,  
'slope': 0.4505162810238786,  
'ca': -0.6360989756865822,  
'thal': -0.67663372723561}
```

```
#Visualize feature importance
feature_df = pd.DataFrame(feature_dict, index=[0])
feature_df.T.plot.bar(title="Feature Importance", legend=False)
```



```
pd.crosstab(df["sex"], df["target"])
```

```
pd.crosstab(df["slope"], df["target"])
```

```
In [89]: pd.crosstab(df["sex"], df["target"])
```

```
Out[89]:
```

	target	0	1
sex			
0	24	72	
1	114	93	

```
In [90]: pd.crosstab(df["slope"], df["target"])
```

```
Out[90]:
```

	target	0	1
slope			
0	12	9	
1	91	49	
2	35	107	

## 6. Experimentation

If you haven't hit your evaluation metric yet ... ask yourself...

- Could you collect more data?
- Could you try a better model? Like CatBoost or XGBoost?
- Could you improve the current models? (beyond what we're done so far)
- If your model is good enough (you have hit your evaluation metric) how would you export it and share it with other?