

ML-10

NEURAL NETWORKS DEEP LEARNING, TRANSFER LEARNING AND
TENSORFLOW 2

TensorFlow

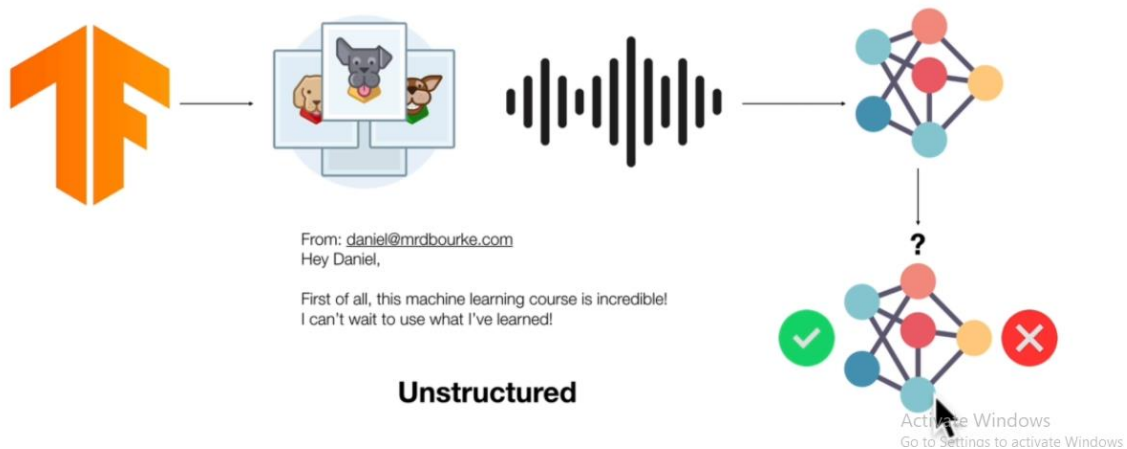
TensorFlow is a deep learning or numerical computing library.

TensorFlow is used to build deep learning and neural network models to gain insights out of unstructured data.

Deep learning is more suited to unstructured data such as photos or audio waves or natural language text. Unstructured data that doesn't really have any specific structure

Now tensor flow allows us to build a deep learning model kind of like a neural network.

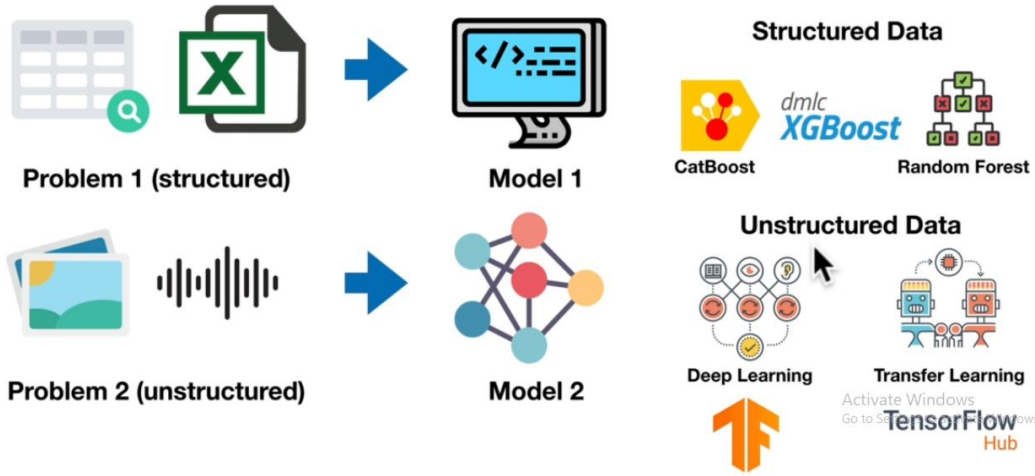
What is TensorFlow?



Why TensorFlow?

- **Write fast deep learning code in Python (able to run on a GPU)**
- **Able to access many pre-built deep learning models**
- **Whole stack: preprocess, model, deploy**
- **Originally designed and used in-house by Google (now open-source)**

Choosing a model (throwback)



Neural Network

Deep Learning

Kind of Deep Learning problems

What kind of deep learning problems are there?



From: daniel@mrdibourke.com
Hey Daniel,


First of all, this machine learning course is incredible!
I can't wait to use what I've learned!

Classification



Hey Siri, where's the nearest cafe?


**Sequence to sequence
(seq2seq)**



Object detection

Activate Windows

What is transfer learning? Why use transfer learning?



- Take what you know in one domain and apply it to another.
- Starting from scratch can be expensive and time consuming.
- Why not take advantage of what's already out there?



Activate Windows
Go to Settings to activate Windows.


Workflow

What are we going to cover?

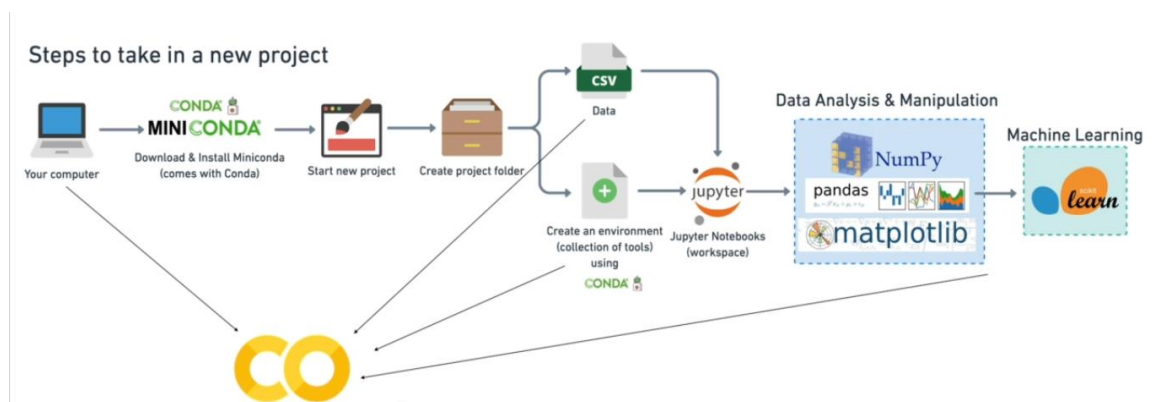
A Scikit-Learn workflow



A TensorFlow workflow

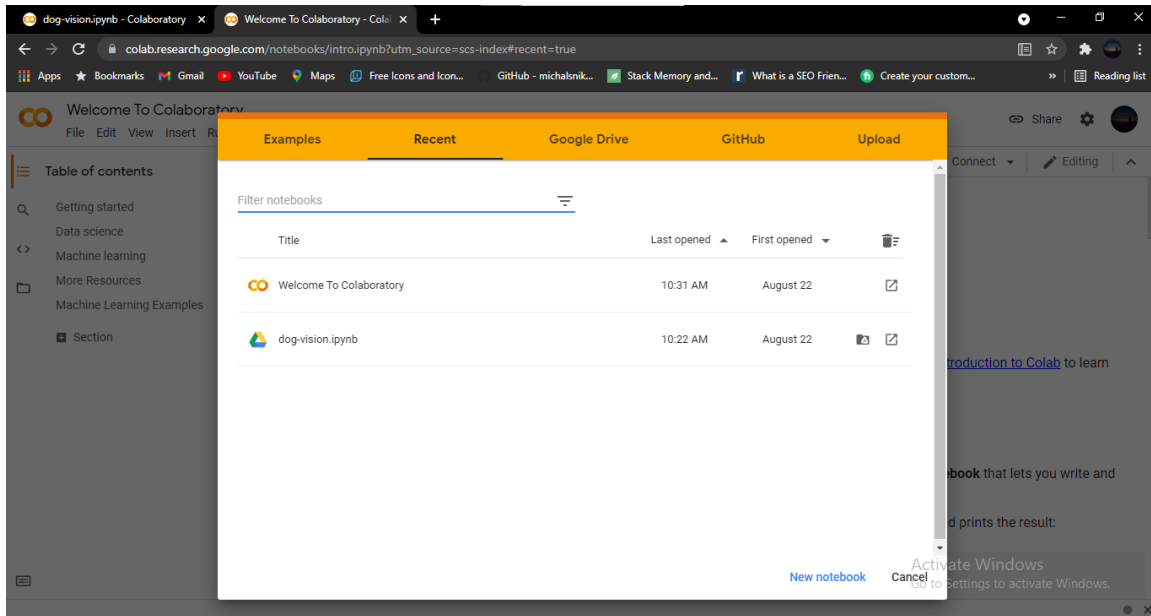


Steps to take in a new project



Setup colab

- Search for colab in google
- Go to colab
- Check for new python notebook



Put the data in google drive and mount google drive

Unzip Files

```
# Unzip the uploaded data into google Drive
# !unzip "drive/MyDrive/Dog-vision/dog-breed-identification.zip" -
d "drive/MyDrive/Dog-vision/"
```

End-to-end Multi-class dog Breed Classification

This notebook builds an end-to-end multi-class image classifier using TensorFlow 2.0 and TensorFlow Hub

1. Problem

Identifying the breed of a dog given an image of a dog.

If I was sitting at that cafe and took a photo of that dog. Could my machine learning model, the one that we're going to build, tell me what breed it is.

2. Data

The data we're using is from Kaggle's dog breed identification competition

3. Evaluation

The evaluation is a file with prediction probabilities for each dog breed of each test image.

4. Features

Some information about the data:

- We're dealing with images (unstructured data) so it's probably best we use deep learning/transfer learning.
- There are 120 breeds of dogs (this means there are 120 different classes).
- There are around 10,000+ images in the training set (these images have labels)
- There are around 10,000+ images in the test set (these images have no labels, because we'll want to predict them)

Get our workspace ready

- Import TensorFlow Hub
- Make sure we're using a GPU

For enable GPU in colab

- Go to **Runtime**
- Click **Change runtime complexity**
- Check for **GPU**
- Click **save**


```
# Import necessary tools
# import tensorflow into colab
import tensorflow as tf
print("TF version", tf.__version__)
# import tensorflow hub
import tensorflow_hub as hub
print("TF Hub version", hub.__version__)

# Check for GPU availability
print("GPU", "available (YESSSSS!!!)" if tf.config.list_physical_device
s("GPU") else "not available :(" )
```



```
# Import necessary tools
# import tensorflow into colab
import tensorflow as tf
print("TF version", tf.__version__)
# import tensorflow hub
import tensorflow_hub as hub
print("TF Hub version", hub.__version__)

# Check for GPU availability
print("GPU", "available (YESSSSS!!!)" if tf.config.list_physical_devices("GPU") else "not available :(" )
```

TF version 2.6.0
TF Hub version 0.12.0
GPU available (YESSSSS!!!)

1. Getting data ready (turning into Tensors)

With all machine learning models, our data has to be in numerical format. So, that's what we'll be doing first. Turning our images into Tensors (numerical representations)

Let's start by accessing our data and checking out the labels.


```
# Checkout the labels of our data
import pandas as pd
label_csv = pd.read_csv("drive/MyDrive/Dog-vision/labels.csv")
print(label_csv.describe())
label_csv.head()
```

```
# Checkout the labels of our data
import pandas as pd
label_csv = pd.read_csv("drive/MyDrive/Dog-vision/labels.csv")
print(label_csv.describe())
label_csv.head()
```


	id	breed
count	10222	10222
unique	10222	120
top	4c8bb929ad8d021729749ef343251fde	scottish_deerhound
freq	1	126

	id	breed
0	000bec180eb18c7604dcecc8fe0dba07	boston_bull
1	001513dfcb2ffafc82ccc4d8bbaba97	dingo
2	001cdf01b096e06d78e9e5112d419397	pekinese
3	00214f311d5d2247d5dfe4fe24b2303d	bluetick
4	0021f9ceb3235effd7fcde7f7538ed62	golden_retriever

```
# How many images are there of each breed?
# label_csv["breed"].value_counts()
label_csv["breed"].value_counts().plot.bar(figsize = (20,10))
```



```
# label_csv["breed"].value_counts().mean()  
label_csv["breed"].value_counts().median()
```



```
# Let's view an image  
from IPython.display import Image  
Image("drive/MyDrive/Dog-  
vision/train/001513dfcb2ffafc82cccf4d8bbaba97.jpg")
```



Getting images and their labels

Let's get a list of all of our image file pathnames

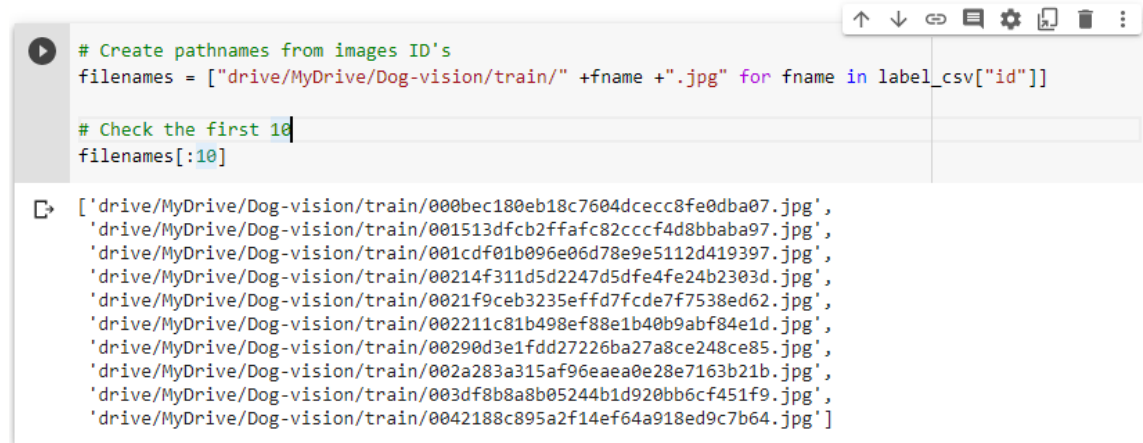
```
# Create pathnames from images ID's
```

```

filenames = ["drive/MyDrive/Dog-
vision/train/" +fname +".jpg" for fname in label_csv["id"]]

# Check the first 10
filenames[:10]

```



```

# Create pathnames from images ID's
filenames = ["drive/MyDrive/Dog-vision/train/" +fname +".jpg" for fname in label_csv["id"]]

# Check the first 10
filenames[:10]

```

```

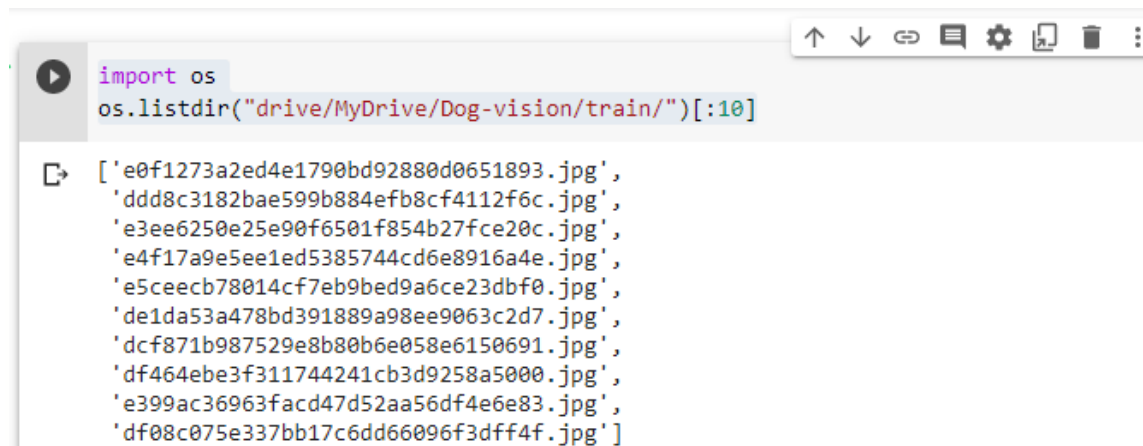
['drive/MyDrive/Dog-vision/train/000bec180eb18c7604dcecc8fe0dba07.jpg',
'drive/MyDrive/Dog-vision/train/001513dfcb2ffafc82cccf4d8bbaba97.jpg',
'drive/MyDrive/Dog-vision/train/001cdf01b096e06d78e9e5112d419397.jpg',
'drive/MyDrive/Dog-vision/train/00214f311d5d2247d5dfe4fe24b2303d.jpg',
'drive/MyDrive/Dog-vision/train/0021f9ceb3235effd7fcde7f7538ed62.jpg',
'drive/MyDrive/Dog-vision/train/002211c81b498ef88e1b40b9abf84e1d.jpg',
'drive/MyDrive/Dog-vision/train/00290d3e1fdd27226ba27a8ce248ce85.jpg',
'drive/MyDrive/Dog-vision/train/002a283a315af96eaea0e28e7163b21b.jpg',
'drive/MyDrive/Dog-vision/train/003df8b8a8b05244b1d920bb6cf451f9.jpg',
'drive/MyDrive/Dog-vision/train/0042188c895a2f14ef64a918ed9c7b64.jpg']

```

```

import os
os.listdir("drive/MyDrive/Dog-vision/train/")[:10]

```



```

import os
os.listdir("drive/MyDrive/Dog-vision/train/")[:10]

```

```

['e0f1273a2ed4e1790bd92880d0651893.jpg',
'ddd8c3182bae599b884efb8cf4112f6c.jpg',
'e3ee6250e25e90f6501f854b27fce20c.jpg',
'e4f17a9e5ee1ed5385744cd6e8916a4e.jpg',
'e5ceecb78014cf7eb9bed9a6ce23dbf0.jpg',
'de1da53a478bd391889a98ee9063c2d7.jpg',
'dcf871b987529e8b80b6e058e6150691.jpg',
'df464ebe3f311744241cb3d9258a5000.jpg',
'e399ac36963facd47d52aa56df4e6e83.jpg',
'df08c075e337bb17c6dd66096f3dff4f.jpg']

```

```

# Check whether number of filenames matches number of actual image files
import os
if len(os.listdir("drive/MyDrive/Dog-
vision/train/")) == len(filenames):
    print("Filenames match actual amount of files!!! Proceed.")
else:
    print("Filenames do not match actual amount of files, check the target directory.")

```


Output : Filenames match actual amount of files!!! Proceed.

```
# One more check / checking image in 9000 index  
Image(filename[9000])
```



```
# Checking breed of dog in upper image  
label_csv["breed"][9000]
```

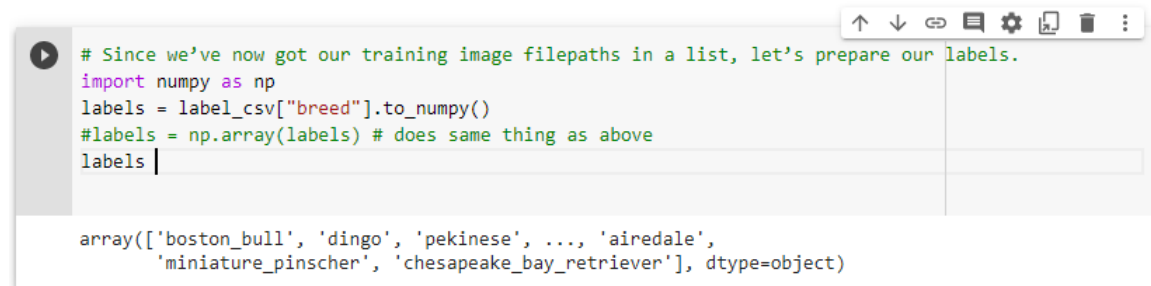
▶ # Checking breed of dog in upper image
label_csv["breed"][9000]

'tibetan_mastiff'

Since we've now got our training image filepaths in a list, let's prepare our labels.

CodeText

```
# Since we've now got our training image filepaths in a list, let's prepare our labels.
import numpy as np
labels = label_csv["breed"].to_numpy()
#labels = np.array(labels) # does same thing as above
labels
```



A Jupyter Notebook cell containing the following code:

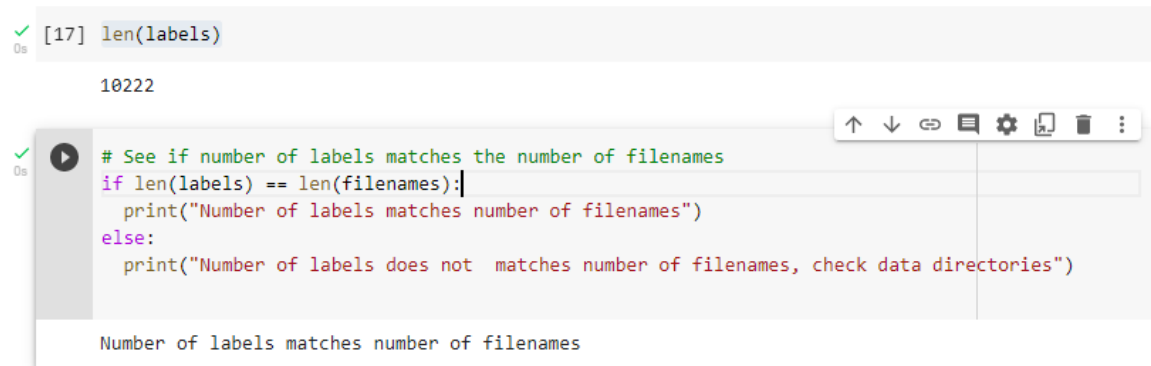
```
# Since we've now got our training image filepaths in a list, let's prepare our labels.
import numpy as np
labels = label_csv["breed"].to_numpy()
#labels = np.array(labels) # does same thing as above
labels
```

The output of the cell is:

```
array(['boston_bull', 'dingo', 'pekinese', ..., 'airedale',
      'miniature_pinscher', 'chesapeake_bay_retriever'], dtype=object)
```

```
len(labels)
```

```
# See if number of labels matches the number of filenames
if len(labels) == len(filenames):
    print("Number of labels matches number of filenames")
else:
    print("Number of labels does not matches number of filenames, check data directories")
```



The first part of the screenshot shows the execution of `len(labels)` in a Jupyter Notebook cell, resulting in the output `10222`.

The second part shows a Jupyter Notebook cell with the following code:

```
# See if number of labels matches the number of filenames
if len(labels) == len(filenames):
    print("Number of labels matches number of filenames")
else:
    print("Number of labels does not matches number of filenames, check data directories")
```

The output of this cell is:

```
Number of labels matches number of filenames
```

```
# Find the unique labels value
unique_breeds = np.unique(labels)
len(unique_breeds)
```

```
# Find the unique labels value
unique_breeds = np.unique(labels)
len(unique_breeds)
```

120

```
# Turn a single label into an array of Booleans
print(labels[0])
labels[0] == unique_breeds
```

```
# Turn a single label into an array of Booleans
print(labels[0])
labels[0] == unique_breeds
```

```
boston_bull
array([False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, True, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False, False, False, False, False, False, False,
       False, False, False])
```

```
# Turn every label into an array of Booleans array
boolean_labels = [label == unique_breeds for label in labels]
boolean_labels[:2]
```

```
# Turn every label into an array of Booleans array
boolean_labels = [label == unique_breeds for label in labels]
boolean_labels[:2]

[array([False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, True, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False]),
 array([False, False, False, False, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
```

```
len(boolean_labels)
```

```
# Example: Turning Boolean array into integer
print(labels[0]) # Original label
print(np.where(unique_breeds == labels[0])) # index where label occurs
print(boolean_labels[0].argmax()) # index where label occurs in boolean
array
print(boolean_labels[0].astype(int)) # there will be 1 where the sample
label occurs
```



```
[25] len(boolean_labels)
```

10222

```
# Example: Turning Boolean array into integer
print(labels[0]) # Original label
print(np.where(unique_breeds == labels[0])) # index where label occurs
print(boolean_labels[0].argmax()) # index where label occurs in boolean array
print(boolean_labels[0].astype(int)) # there will be 1 where the sample label occurs
```

[illegible]

```
print(labels[2])
print(boolean_labels[2].astype(int))
```

```
filenames[:10]
```

```
[28] print(labels[2])
      print(boolean_labels[2].astype(int))
```

[illegible]

```
filenames[:10]
```

```
[ 'drive/MyDrive/Dog-vision/train/000bec180eb18c7604dcecc8fe0dba07.jpg',  
  'drive/MyDrive/Dog-vision/train/001513dfcb2ffaf8c82ccc4d8bbabab97.jpg',  
  'drive/MyDrive/Dog-vision/train/001cdf01b096e06d78e95112d419397.jpg',  
  'drive/MyDrive/Dog-vision/train/00214f311d5d2247d5dfe4fe24b2303d.jpg',  
  'drive/MyDrive/Dog-vision/train/0021f9ceb3235effd7fcde7f7538ed62.jpg',  
  'drive/MyDrive/Dog-vision/train/002211c81b498ef88e1b40b9abf84e1d.jpg',  
  'drive/MyDrive/Dog-vision/train/00290d3e1fdd27226ba27a8ce248ce85.jpg',  
  'drive/MyDrive/Dog-vision/train/002a283a315af96eaea0e28e7163b21b.jpg',  
  'drive/MyDrive/Dog-vision/train/003df8b8a8b05244b1d920bb6cf451f9.jpg',  
  'drive/MyDrive/Dog-vision/train/0042188c895a2f14ef64a918ed9c7b64.jpg']
```

Creating our own validation set

Since the dataset from Kaggle doesn't come with a validation set, we're going to create our own.

OS

```
# setup x and y variables

x = filenames
y = boolean_labels
```

We're going to start off experimenting with ~1000 images and increase as needed

```
# Set number of images to use for experimenting
NUM_IMAGES = 1000 #@param {type:"slider", min:1000,max:10000,step:100}
```

```
# Let's split our data into train and validation sets
from sklearn.model_selection import train_test_split

# Split them into training and validation to total size NUM_IMAGES
X_train, X_val, y_train, y_val = train_test_split(x[:NUM_IMAGES],
                                                  y[:NUM_IMAGES],
                                                  test_size=0.2,
                                                  random_state=42)

len(X_train), len(y_train), len(X_val), len(y_val)
```



```
# Let's have a peek at the training
X_train[:5], y_train[:2]
```

```
# Let's have a geez at the training
X_train[:5], y_train[:2]

(['drive/MyDrive/Dog-vision/train/00bee065dcec471f26394855c5c2f3de.jpg',
 'drive/MyDrive/Dog-vision/train/0d2f9e12a2611d911d91a339074c8154.jpg',
 'drive/MyDrive/Dog-vision/train/1108e48ce3e2d7d7fb527ae6e40ab486.jpg',
 'drive/MyDrive/Dog-vision/train/0dc3196b4213a2733d7f4bdcd41699d3.jpg',
 'drive/MyDrive/Dog-vision/train/146fbfac6b5b1f0de83a5d0c1b473377.jpg'],
 [array([False, False, False, False, False, False, False, False, False,
         False, False, False, False, False, False, False, False, False,
         False, False, False, False, False, False, False, False, True,
         False, False, False, False, False, False, False, False, False,
         False, False, False, False, False, False, False, False, False,
         False, False, False, False, False, False, False, False, False,
         False, False, False, False, False, False, False, False, False,
         False, False, False, False, False, False, False, False, False,
         False, False, False, False, False, False, False, False, False,
         False, False, False, False, False, False, False, False, False,
         False, False, False, False, False, False, False, False, False,
         False, False, False]),
  array([False, False, False, False, False, False, False, False, False,
         False, False, False, False, False, False, False, False, False,
         False, False, False, False, False, False, False, False, False,
         False, False, False, False, False, False, False, False, False,
         False, False, False, False, False, False, False, False, False,
         False, False, False, False, False, False, False, False, False,
         False, False, False])])
```

Preprocessing Images (turning images into Tensors)

To preprocess our images into Tensors we're going to write a function which does a few things:

2. Take an image filepath as input
3. Use TensorFlow to read the file and save it to a variable, image
4. turn our image (a jpg) into Tensors
5. Resize the image to be a shape of (224,224)
6. Return the modified image

Before we do, let's see what importing an image looks likes.

```
# convert image to Numpy array
from matplotlib.pyplot import imread
image = imread(filename[42])
image.shape
```

```
# convert image to Numpy array
from matplotlib.pyplot import imread
image = imread(filenamees[42])
image.shape
```

(257, 350, 3)

```
image
array([[ 89, 137,  87],
       [ 76, 124,  74],
       [ 63, 111,  59],
       ...,
       [ 76, 134,  86],
       [ 76, 134,  86],
       [ 76, 134,  86]],

      [[ 72, 119,  73],
       [ 67, 114,  68],
       [ 63, 111,  63],
       ...,
       [ 75, 131,  84],
       [ 74, 132,  84],
       [ 74, 131,  86]],

      [[ 56, 104,  66],
       [ 58, 106,  66],
       [ 64, 112,  72],
       ...,
       [ 71, 127,  82],
       [ 73, 129,  84],
       ...,
       ...,
       ...])
```

```
# every image is between 0 to 255 (RGB pixel values)
image.max(), image.min()
```

```
#turn image into tensors
tf.constant(image)[:2]
```

```
# every image is between 0 to 255 (RGB pixel values)
image.max(), image.min()

(255, 0)

[42] #turn image into tensors
tf.constant(image)[:2]

<tf.Tensor: shape=(2, 350, 3), dtype=uint8, numpy=
array([[ [ 89, 137,  87],
        [ 76, 124,  74],
        [ 63, 111,  59],
        ...,
        [ 76, 134,  86],
        [ 76, 134,  86],
        [ 76, 134,  86]],
       [[ [ 72, 119,  73],
        [ 67, 114,  68],
        [ 63, 111,  63],
        ...,
        [ 75, 131,  84],
        [ 74, 132,  84],
        [ 74, 131,  86]]], dtype=uint8)>
```

Now we've seen what an image looks like as a tensor, let's make a function to pre process them.

1. Take an image filepath as input
2. Use TensorFlow to read the file and save it to a variable, image
3. turn our image (a jpg) into Tensors
4. normalize our image (convert color channel values from 0 to 255 to 0 to 1)
5. Resize the image to be a shape of (224,224)
6. Return the modified image

```
# Define image size
IMG_SIZE = 224

# Create a function for preprocessing image

def process_image(image_path, img_size=IMG_SIZE):
    """
    Takes an image filepath and turns the image into a Tensor.
```

```

"""
#Read in an image file
image = tf.io.read_file(image_path)

# Turn the jpeg image into numerical Tensor with 3 color channels (RGB)
image = tf.image.decode_jpeg(image, channels=3)

# Convert the color channel values from 0-255 to 0-1 values
image = tf.image.convert_image_dtype(image, tf.float32)

# Resize the image to our desired value (224, 224)
image = tf.image.resize(image, size=(IMG_SIZE, IMG_SIZE))

return image

```

Turning our data into batches

Why turn our data into batches?

Let's say you're trying to process 10,000+ images in one go... they all might not fit into memory.

so that's why we do about 32 (this is batch size) images at a time (you can manually adjust the batch size if need be)

In order to use TensorFlow effectively, we need our data in the form of Tensor tuples which look like this: (image, label)

```

## Create a simple function to return a tuple `(image, label)`

def get_image_label(image_path, label):
    """
    Takes an image file path name and the associated label,
    processes the image and returns a tuple of image (image, label).
    """

    image = process_image(image_path)
    return image, label

```

```

# Demo of the above
(process_image(x[42]), tf.constant(y[42]))

```

```
# Demo of the above
(process_image(x[42]), tf.constant(y[42]))

[0.29721776, 0.52466875, 0.33030328],
[0.2948505 , 0.5223015 , 0.33406618]],

[[0.25903144, 0.4537807 , 0.27294815],
 [0.24375686, 0.4407019 , 0.2554778 ],
 [0.2838985 , 0.47213382, 0.28298813],
 ...,
 [0.2785345 , 0.5027992 , 0.31004712],
 [0.28428748, 0.5108719 , 0.32523635],
 [0.28821915, 0.5148036 , 0.32916805]],

[[0.20941195, 0.40692952, 0.25792548],
 [0.24045378, 0.43900946, 0.2868911 ],
 [0.29001117, 0.47937486, 0.32247734],
```

Now we've got a way to turn our data into tuples of Tensors in the form: (image, label), let's make a function to turn all of our data (x and y) into batches!

```
# Define the batch size, 32 is a good start
BATCH_SIZE = 32

# Create a function to turn data into batches
def create_data_batches(x, y=None, batch_size=BATCH_SIZE, valid_data= F
    else, test_data=False): """
    """
    Creates batches of data out of image (x) and label (y) pairs.
    Shuffles the data if it's training data but doesn't shuffle if it's vali
    dation data.
    Also accepts test data as input (no labels).

    """
    # If the data is a test dataset, we probably don't have labels
    if test_data:
        print("Creating test data batches...")
        data = tf.data.Dataset.from_tensor_slices(tf.constant(x)) # only fi
        lepaths (no labels)
        data_batch = data.map(process_image).batch(BATCH_SIZE)
        return data_batch
    # If the data is a valid dataset, we don't need to shuffle it
    elif valid_data: #####
        print("Creating validation data batches...")
        data = tf.data.Dataset.from_tensor_slices((tf.constant(x), # filepa
        ths
                                                    tf.constant(y))) # labels
```

```

    data_batch = data.map(get_image_label).batch(BATCH_SIZE)
    return data_batch

else:
    print("Creating training data batches...")
    # Turn filepaths and labels into Tensors
    data = tf.data.Dataset.from_tensor_slices((tf.constant(x),
                                              tf.constant(y)))

    #Shuffling pathnames and labels before mapping image processor function
    # is faster then shuffling images
    data = data.shuffle(buffer_size = len(x))

    # Create (image, label) tuples (this also turns the image path into
    # a preprocessed image)
    data = data.map(get_image_label)

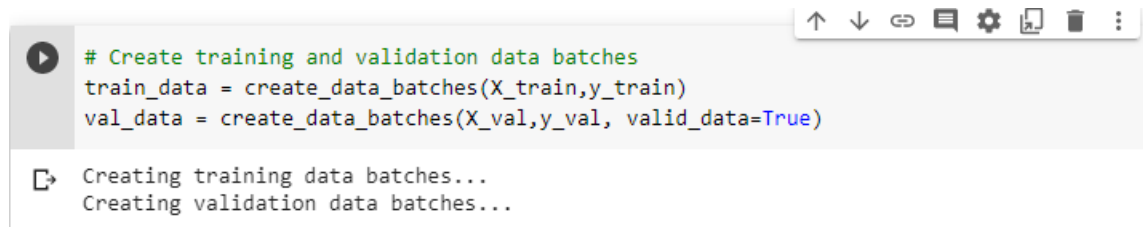
    # Turn the training data into batches
    data_batch = data.batch(BATCH_SIZE)
    return data_batch

```

```

# Create training and validation data batches
train_data = create_data_batches(X_train,y_train)
val_data = create_data_batches(X_val,y_val, valid_data=True)

```



The image shows a Jupyter Notebook interface. At the top, there is a toolbar with icons for undo, redo, insert cell, delete cell, run cell, insert comment, settings, copy, paste, and a menu. Below the toolbar, a code cell contains the following text: `# Create training and validation data batches`, `train_data = create_data_batches(X_train,y_train)`, and `val_data = create_data_batches(X_val,y_val, valid_data=True)`. Below the code cell, an output cell displays the text: `Creating training data batches...` and `Creating validation data batches...`.

```

# # Check out the different attributes of our data batches
train_data.element_spec, val_data.element_spec

```



```
# # Check out the different attributes of our data batches
train_data.element_spec, val_data.element_spec

((TensorSpec(shape=(None, 224, 224, 3), dtype=tf.float32, name=None),
  TensorSpec(shape=(None, 120), dtype=tf.bool, name=None)),
 (TensorSpec(shape=(None, 224, 224, 3), dtype=tf.float32, name=None),
  TensorSpec(shape=(None, 120), dtype=tf.bool, name=None)))
```

Visualizing data batches

our data in batches can be a bit of a hard concept to understand. Let's build a function which helps us visualize what's going on under the hood.

```
import matplotlib.pyplot as plt

# Create a function for viewing images in a data batch

def show_25_images(images, labels):
    """
    Displays 25 images from a data batch.
    """
    # Setup the figure

    plt.figure(figsize=(10,10))

    # Loop through 25 (for displaying 25 images)

    for i in range(25):
        # Create subplots (5 rows, 5 columns)
        ax = plt.subplot(5, 5, i+1)
        # Display an image
        plt.imshow(images[i])
        # Add the image label as the title
        plt.title(unique_breeds[labels[i].argmax()])
        # Turn the grid line off
        plt.axis("off")
```

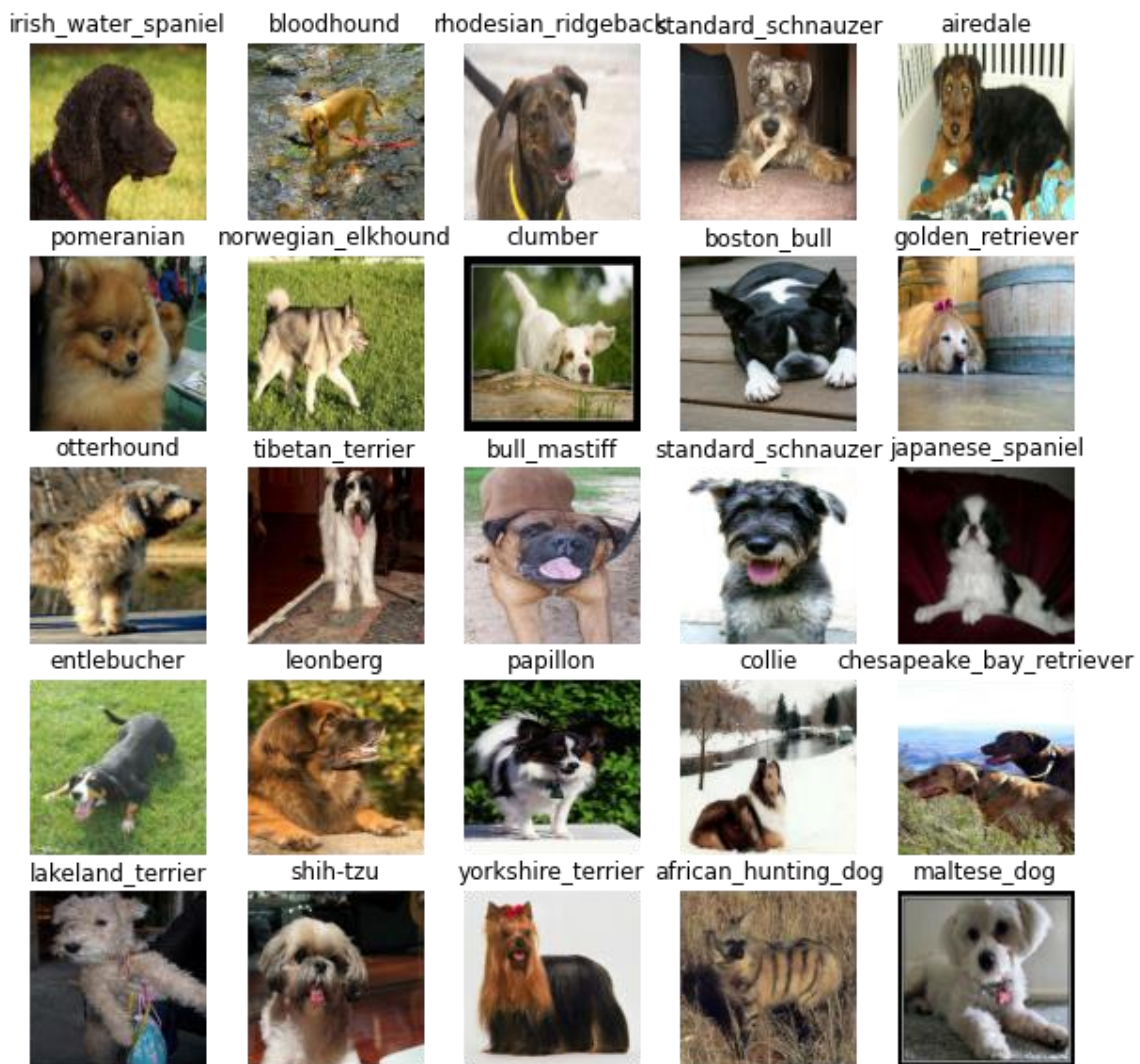
```
#train_images, train_labels = next(train_data.as_numpy_iterator())
#len(train_images), len(train_labels)
```

```
#train_images, train_labels = next(train_data.as_numpy_iterator())
#len(train_images), len(train_labels)
```

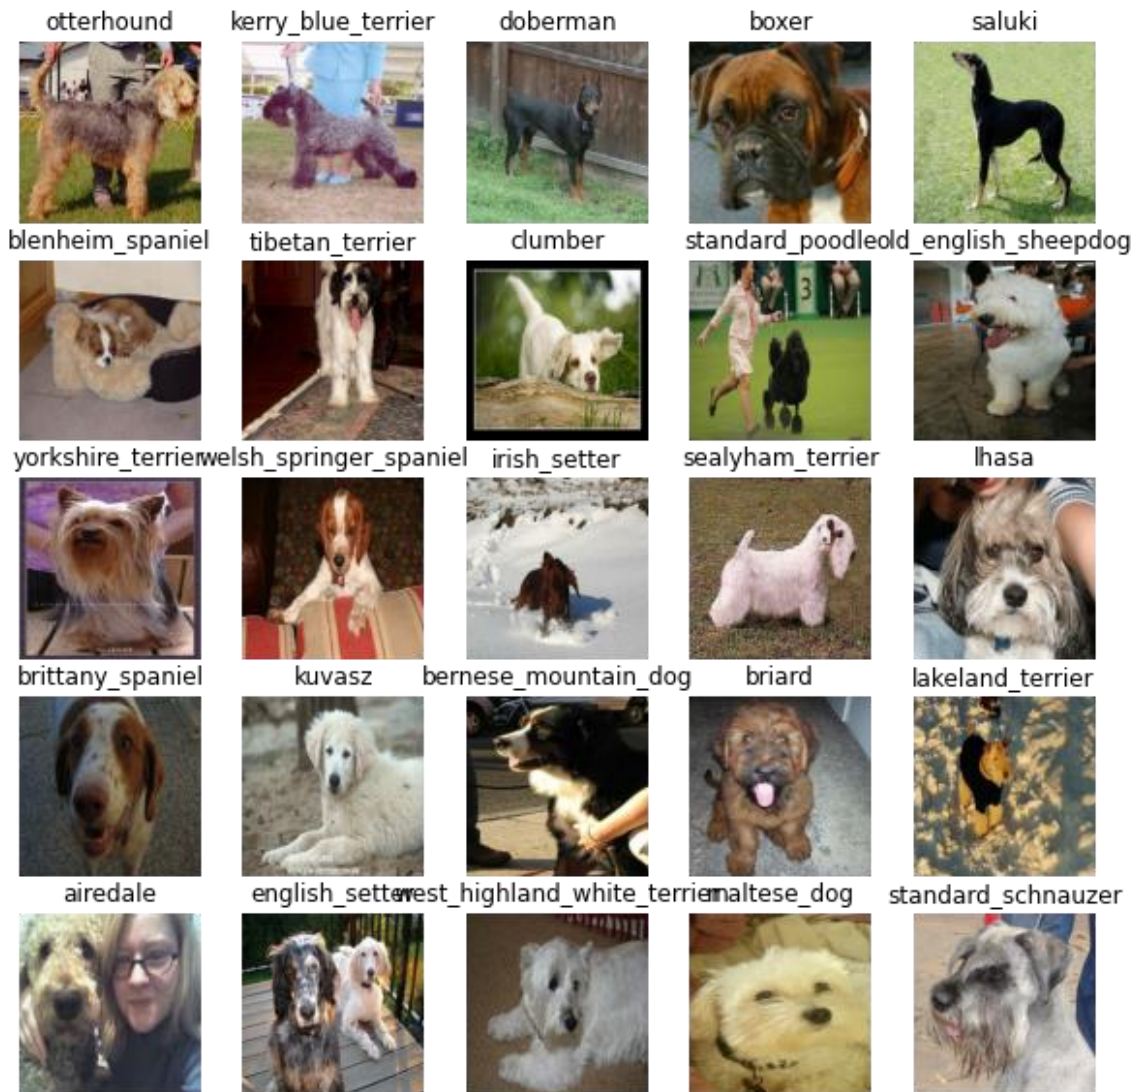
(32, 32)

+ Code + Text

```
# Now let's visualize the data in a training batch
train_images, train_labels = next(train_data.as_numpy_iterator())
show_25_images(train_images, train_labels)
```

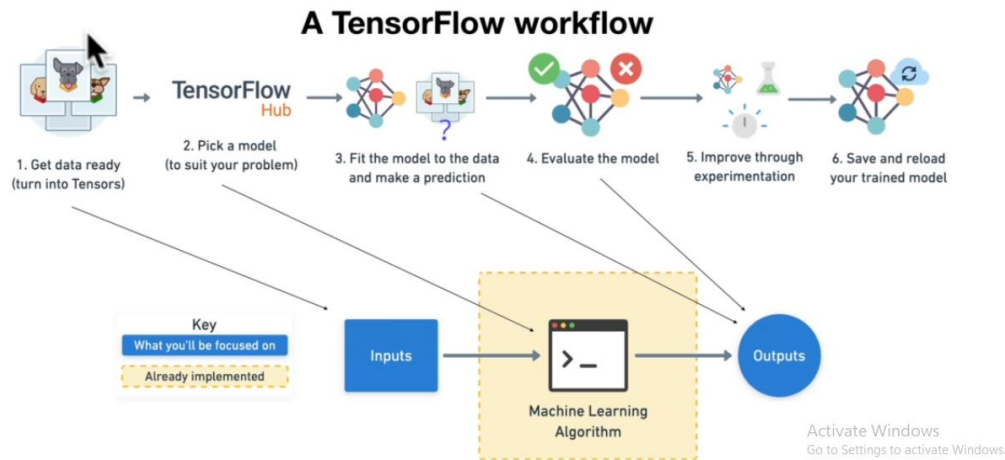


```
# Now let's Visualize our validation set
val_images, val_labels = next(train_data.as_numpy_iterator())
show_25_images(val_images, val_labels)
```



Preparing Our Inputs and Outputs

What we're focused on



Creating and training a model

Now our data is ready, let's prepare it modelling. We'll use an existing model from [TensorFlow Hub](#).

TensorFlow Hub is a resource where you can find pretrained machine learning models for the problem you're working on.
Using a pretrained machine learning model is often referred to as **transfer learning**.

Why use a pretrained model?

Building a machine learning model and training it on lots from scratch can be expensive and time consuming.

Transfer learning helps eliviate some of these by taking what another model has learned and using that information with your own problem.

How do we choose a model?

Since we know our problem is image classification (classifying different dog breeds), we can navigate the [TensorFlow Hub page by our problem domain \(image\)](#).

We start by choosing the image problem domain, and then can filter it down by subdomains, in our case, [image classification](#).

Doing this gives a list of different pretrained models we can apply to our task.

Clicking on one gives us information about the model as well as instructions for using it. For example, clicking on the [mobilenet_v2_130_224](#) model, tells us this model takes an input of images in the shape 224, 224. It also says the model has been trained in the domain of image classification.

Let's try it out.

Building a model

Before we build a model, there are a few things we need to define:

- The input shape (images, in the form of Tensors) to our model.
- The output shape (image labels, in the form of Tensors) of our model.
- The URL of the model we want to use.

These things will be standard practice with whatever machine learning model you use. And because we're using TensorFlow, everything will be in the form of Tensors.

```
# Setup input shape to the model
INPUT_SHAPE = [None, IMG_SIZE, IMG_SIZE, 3] # batch, height, width, col
our channels

# Setup output shape to the model
OUTPUT_SHAPE = len(unique_breeds)
# Setup model URL from TensorFlow Hub
MODEL_URL = "https://tfhub.dev/google/imagenet/mobilenet_v2_130_224/cla
ssification/4"
```

Now we've got the inputs, outputs and model we're using ready to go. We can start to put them together into a **Keras** deep learning model

There are many ways of building a model in TensorFlow but one of the best ways to get started is to use the Keras API.

Defining a deep learning model in Keras can be as straightforward as saying, "here are the layers of the model, the input shape and the output shape, let's go!"

Knowing this, let's create a function which:

- Takes the input shape, output shape and the model we've chosen's URL as parameters.
- Defines the layers in a Keras model in a sequential fashion (do this first, then this, then that).
- Compiles the model (says how it should be evaluated and improved).
- Builds the model (tells it what kind of input shape it'll be getting).
- Returns the model.

All of these steps can be found here: www.tensorflow.org/guide/keras/overview

We'll take a look at the code first, then dicuss each part.

```
# Create a function which builds a Keras model

def create_model(shape=INPUT_SHAPE, output_shape=OUTPUT_SHAPE, model_url=MODEL_URL):
    print("Building model with:", MODEL_URL)

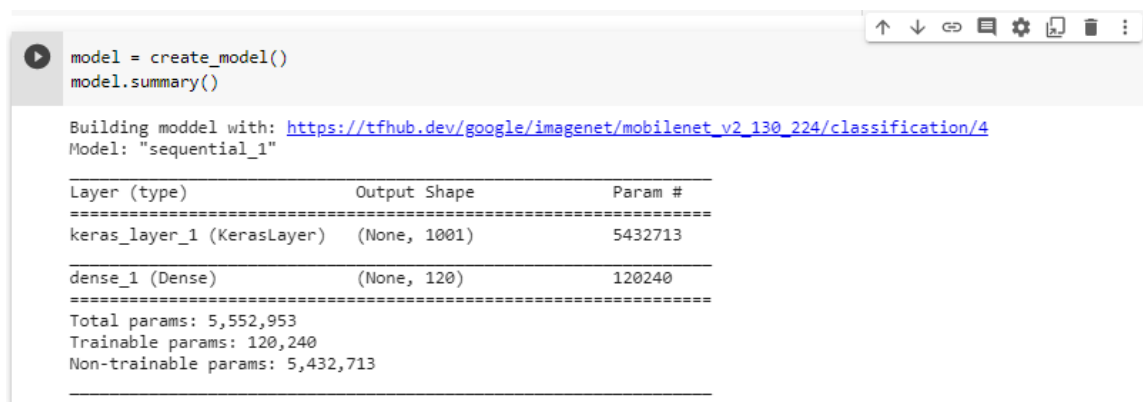
    # Setup the model layers
    model = tf.keras.Sequential([
        hub.KerasLayer(MODEL_URL), #Layer 1 (input layer)
        tf.keras.layers.Dense(units=OUTPUT_SHAPE, activation="softmax") #Layer 2 (output layer)
    ])

    # Compile the model
    model.compile(
        loss = tf.keras.losses.CategoricalCrossentropy(),
        optimizer = tf.keras.optimizers.Adam(),
        metrics = ["accuracy"]
    )

    # Build the model
    model.build(INPUT_SHAPE)

    return model
```

```
model = create_model()
model.summary()
```



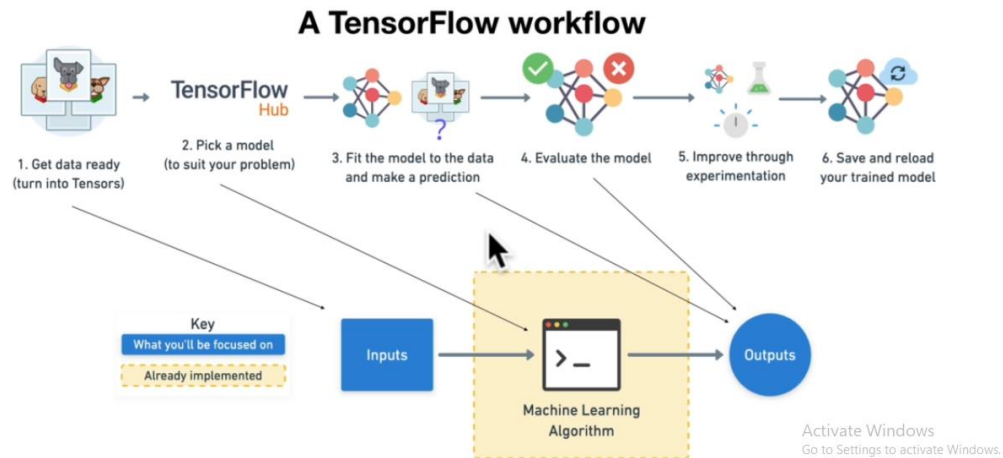
```
model = create_model()
model.summary()
```

Building model with: https://tfhub.dev/google/imagenet/mobilenet_v2_130_224/classification/4
Model: "sequential_1"

Layer (type)	Output Shape	Param #
keras_layer_1 (KerasLayer)	(None, 1001)	5432713
dense_1 (Dense)	(None, 120)	120240

Total params: 5,552,953
Trainable params: 120,240
Non-trainable params: 5,432,713

What we're focused on



Which activation? Which loss?

Binary classification	Multi-class classification
Activation: Sigmoid	Activation: Softmax
Loss: Binary Crossentropy	Loss: Categorical Crossentropy

Activate Windows
Go to Settings to activate Windows.